

Joins in BigQuery


INTRODUCTION TO BIGQUERY



Matthew Forrest
Field CTO

Joining data in BigQuery

- **INNER JOIN** : Values exist in both tables.
- **LEFT JOIN** : **All rows in left table**, matches right table.
- **RIGHT JOIN** : **All rows in right table**, matches left table.
- **FULL JOIN** : All rows from **both tables**, matches **and non matches**.
- **CROSS JOIN** : **Every row** matched to **every row** from both tables.



Real life examples of joins

Customers: left table

Orders: right table

- **INNER JOIN** : Matching customers and their orders
- **LEFT JOIN** : Showing all customers, even if they haven't placed any orders
- **RIGHT JOIN** : Showing all orders, even if there is a missing customer ID
- **FULL JOIN** : Showing all customers and all orders, even if some haven't interacted
- **CROSS JOIN** : Match every order to every customer with no conditions

INNER JOIN

- Only returns matching results from both datasets

```
SELECT
```

```
  c.customer_id, s.product_name
```

```
FROM customers c
```

```
-- The INNER keyword is optional
```

```
JOIN sales_data s
```

```
ON c.customer_id = s.customer_id;
```

```
| customer_id | product_name          |
|-----|-----|
| 1          | Bluetooth Headphones |
| 2          | Running Shoes        |
```

LEFT JOIN

- Returns all rows from the LEFT dataset

SELECT

c.customer_id, s.product_name

FROM customers c

LEFT JOIN sales_data s

ON c.customer_id = s.customer_id;

customer_id	product_name
1	Bluetooth Headphones
2	Running Shoes
3	null

RIGHT JOIN

- Returns all rows from the RIGHT dataset

SELECT

c.customer_id, s.product_name

FROM customers c

RIGHT JOIN sales_data s

ON c.customer_id = s.customer_id;

customer_id	product_name
1	Bluetooth Headphones
2	Running Shoes
null	External Microphone

OUTER JOIN

- A "RIGHT-LEFT" join: all rows from both RIGHT a and LEFT datasets

```
SELECT
```

```
  c.customer_id, s.product_name
```

```
FROM customers c
```

```
OUTER JOIN sales_data s
```

```
ON c.customer_id = s.customer_id;
```

customer_id	product_name
1	Bluetooth Headphones
2	Running Shoes
3	null
null	External Microphone

SELF or CROSS JOIN

- A cartesian join - every row with every row

SELECT

```
c.customer_id,  
s.product_name,
```

```
-- Adding table names separated  
-- by a comma is a CROSS JOIN  
-- Order is determined by the  
-- left table, here "customers"
```

FROM customers c, sales_data s;

customer_id	product_name
1	Bluetooth Headphones
1	null
2	Bluetooth Headphones
2	null
3	null
3	Bluetooth Headphones

Joins and UNNEST

- Also used to join unnested data

SELECT

```
c.customer_id,  
payments.method
```

FROM customers c,

UNNEST(

```
customers.payment_methods
```

) payments;

customer_id	product_name
1	Visa
1	Mastercard
1	Venmo
1	Paypal
2	Amex
2	Visa

Let's practice!

INTRODUCTION TO BIGQUERY

Data manipulation language (DML) statements

INTRODUCTION TO BIGQUERY



Matt Forrest
Field CTO

Overview of data manipulation in BigQuery

- `INSERT` : Add new rows of data.
- `UPDATE` : Modify existing values in a row.
- `DELETE` : Remove unwanted data from a tables.
- `MERGE` : Statement that can combine `INSERT` , `UPDATE` , and `DELETE` statements into one statement.
- `CREATE TABLE AS` : Creates a new table from a query result.

Considerations and performance

- Group DML statements together when possible rather than running them individually
- You must use a `WHERE` condition when running an `UPDATE` statement
- Consider using table partitions and clusters

¹ <https://cloud.google.com/bigquery/docs/reference/standard-sql/data-manipulation-language>

INSERT

- Add records to tables

-- Define the columns in the parentheses

```
INSERT INTO customers (customer_id, name, email)
```

-- Each value is a row to be inserted

```
VALUES (1, "John Doe", "john.doe@example.com"),  
(2, "Jane Doe", "jane.doe@example.com"),  
(3, "Alice Smith", "alice.smith@example.com");
```

UPDATE

- Changing data based on a condition

```
UPDATE customers
-- Set one column for each SET statement
SET email = "john.doe@newdomain.com"
-- Make sure to include where otherwise all
-- rows will be updated
WHERE customer_id = 1;
```

- UPDATE together with subqueries or joins

```
UPDATE customers c
SET c.email = e.email
FROM emails e
WHERE c.customer_id = 1;
```

DELETE

- `DELETE` permanently removes records and can't be reversed

```
DELETE FROM customers
```

```
-- Include WHERE to ensure only specific rows are deleted
```

```
WHERE customer_id = 3;
```

```
DELETE FROM customers c
```

```
JOIN emails e USING (customer_id)
```

```
WHERE email = 'john.doe@newdomain.com'
```


MERGE

- Combines `INSERT` , `UPDATE` , and `DELETE` in a single operation

```
-- Sets the customers table as the target
MERGE customers AS target
-- The source is set to new_customers
USING new_customers AS source
-- Matching condition
ON target.customer_id = source.customer_id
-- If the emails do not match, update the email
WHEN MATCHED AND target.email != source.email THEN
    UPDATE SET email = source.email
-- If the match is not met, insert the record
WHEN NOT MATCHED THEN
    INSERT (customer_id, name, email) VALUES
    (source.customer_id, source.name, source.email);
```

CREATE TABLE

- Create new tables from queries

```
CREATE TABLE active_customers AS
SELECT customer_id, name, email
FROM customers
WHERE last_active_date > DATE_SUB(CURRENT_DATE(), INTERVAL 30 DAY);
```

Let's practice!

INTRODUCTION TO BIGQUERY

Query optimization strategies

INTRODUCTION TO BIGQUERY



Matt Forrest
Field CTO

Three rules of thumb

There are three main optimization rules:

1. **Reduce** the amount of data that needs to be **processed**
2. **Optimize** the query operations
3. **Reduce** the output size of your query

¹ <https://cloud.google.com/bigquery/docs/best-practices-performance-compute#use-bi-engine>

Reducing the amount of data

- Avoid using `SELECT *`, and only select the columns of data we need
- Filter the amount of data in our CTEs early and often
- Filter our data using the `WHERE` clause early and often

Optimizing joins

- Make sure we reduce the data we need using CTEs.
- Join using an `INT64` data type.

```
WITH filter_my_data AS (SELECT
-- Filter data with
-- WHERE in the CTE first
)
SELECT
-- This query will run
-- faster with less data
JOIN a USING (user_id)
```

Optimizing the WHERE clause

- In BigQuery, use
 - `BOOL`
 - `INT`
 - `FLOAT`
 - `DATE`
- Data types with `WHERE` , `STRING` , or `BYTE` are **not optimal**.

Not optimal

```
SELECT user_id, date_ordered  
FROM dataset.table  
WHERE product = 'shoes'
```

Optimal

```
SELECT user_id, date_ordered  
FROM dataset.table  
WHERE product_id = 1234
```


ORDER BY optimizations

- `ORDER BY` should always be at the outermost (end) of our query
- The only exception to this is using `ORDER BY` within a window clause

ORDER BY without optimization

Not optimal

```
WITH order_total AS (SELECT
  user_id,
  sum(product_price) as order_sum
FROM orders
GROUP BY user_id
-- Order by is not at the end of the query
ORDER BY last_purchase_date
)
SELECT order_total.order_sum,
  users.user_name
FROM dataset.users users
JOIN order_total USING (user_id);
```

ORDER BY with optimization

Optimal

```
WITH order_total AS (SELECT
  user_id,
  last_purchase_date
  sum(product_price) as order_sum
GROUP BY user_id
)
SELECT order_total.order_sum,
  users.user_name
FROM dataset.users users
JOIN a USING (user_id)
-- Order by should always be at the end
ORDER BY order_total.last_purchase_date;
```

Using EXISTS vs. COUNT

- If we only need to know if a record is in the table, using `EXISTS`
- Avoid using `COUNT` to solve this use case

```
SELECT EXISTS (  
  
    -- Write the main query as a subquery v  
  
    SELECT  
        user_id  
  
    FROM  
        dataset.table  
  
    WHERE  
        product_category = 'home_goods'  
        AND status = 'Closed Account'  
  
);
```

Other optimization methods

- Use approximate aggregate functions such as `APPROX_TOP_SUM` or `APPROX_COUNT_DISTINCT` .
- Many BigQuery tables are partitioned by date - include dates in the `WHERE` clause.

Let's practice!

INTRODUCTION TO BIGQUERY

Congratulations!

INTRODUCTION TO BIGQUERY



Matt Forrest
Field CTO

What we covered

1. BigQuery architecture, background, and comparisons
2. Data ingestion, data types, and unstructured data
3. Querying data in BigQuery with CTEs, aggregations, and WINDOWS
4. Joins, query optimizations, and data manipulation

What's next?

All BigQuery code samples

This page contains code samples for BigQuery. To search and filter code samples for other Google Cloud products, see the [Google Cloud sample browser](#).

FILTER BY

Choose languages

Select all

☐ Java

☐ Python

☐ Go

More

Choose APIs

Select all

☐ BigQuery Storage

☐ BigQuery Connection API

Filter results

Get configuration metadata

Get transfer configuration metadata.

BigQuery Google Cloud

BigQuery Data Transfer Service Java

Get dataset properties

Retrieve the properties of a dataset.

Google Cloud BigQuery Python

Node JS JavaScript Go Java

Enable large results

Query enables large result sets using legacy SQL.

BigQuery Google Cloud Python

Node JS JavaScript Java Go

Query Sheets with a temporary table

Cancel a job

Attempt to cancel a job.

Table exists

A function to check whether a table exists

Congrats!

INTRODUCTION TO BIGQUERY